

Kubernetes from Basic to Advanced



kubernetes

Session #02

Introduction to Kubernetes



kubernetes

Session Contents



- Kubernetes Clusters: OnPrem, Cloud, Local
- CLI Tooling
- Declarative Configuration
- Pods
- Namespace

Kubernetes Clusters: OnPrem, Cloud, Local



Kubernetes Clusters



- Set of nodes running Kubernetes agents
- Control Plane (Master nodes) and Nodes (Worker Nodes)
- High availability achieved running 1+ Control Plane Nodes
- Needs to have an external load balancer to achieve HA
- May have different nodes with different hardware
- Several layers of configuration and maintenance
 - VMs / Bare Metal Machines
 - Guest OS
 - Kubernetes software

Kubernetes Flavours



- Kubernetes Vanilla means no add-ons
- Several Flavours with different focus on added services
 - Suse Rancher
 - Red Hat OpenShift
 - VMWare Tanzu
 - Mirantis Kubernetes Engine, ...
- All flavours add services on top of Kubernetes Vanilla meaning if you only use vanilla resources you are vendor-free

Kubernetes Clusters: OnPrem



- All node types needs to be created and maintain manually
- For Kubernetes Vanilla, all installation process must use [kubeadm](#) tool
- Requirements
 - One or more machines running a deb/rpm-compatible Linux OS; for example: Ubuntu or CentOS.
 - 2 GiB or more of RAM per machine--any less leaves little room for your apps.
 - At least 2 CPUs on the machine that you use as a control-plane node.
 - Full network connectivity among all machines in the cluster. You can use either a public or a private network.
- Control Plane nodes needs to be on Linux

Kubernetes Clusters: OnPrem



- Benefits
 - Full control on hardware and software
 - Allow making some changes on specific components (etcd, etc.)
- Challenges
 - Node autoscaling
 - External communication
 - Control Plane HA
 - Kubernetes version update

Kubernetes Clusters: Cloud



- Cloud Providers have their own flavours
 - Azure Kubernetes Service (AKS)
 - Google Kubernetes Engine (GKE)
 - Amazon Elastic Kubernetes Service (EKS)
 - DigitalOcean Kubernetes
- All are compatible with Kubernetes Vanilla
- Additional services may cause vendor-lock (like Azure AD integration)
 - Identity
 - Policies
 - Secrets Management, ...

Kubernetes Clusters: Cloud



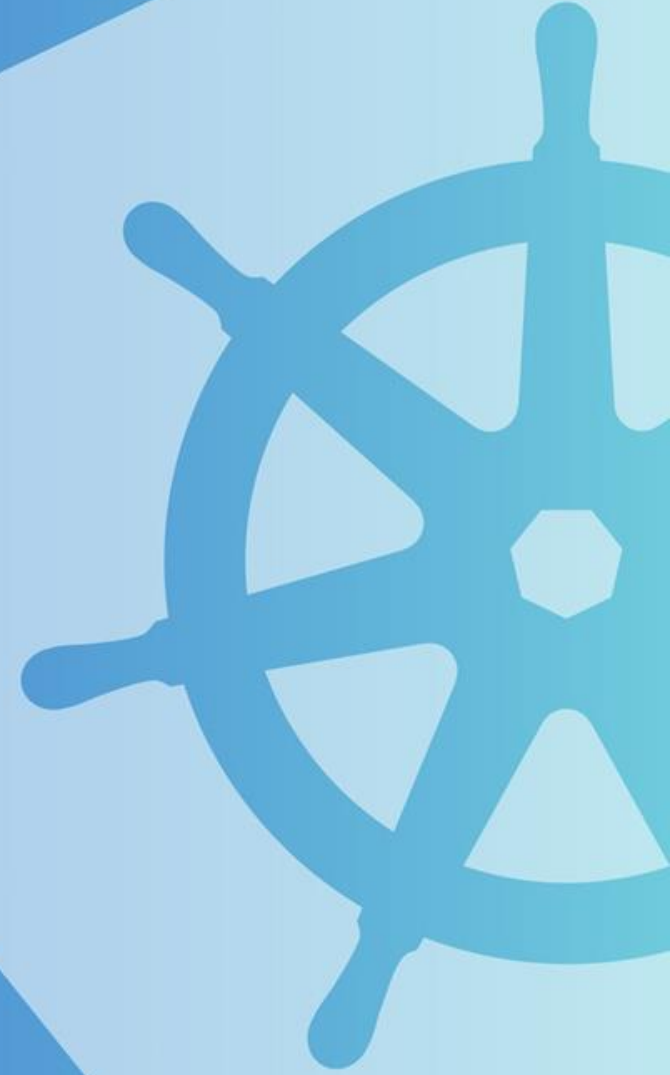
- Kubernetes as a Service
- Cloud Providers uses the concept of Managed Services
- On all of them, Control Plane nodes are totally managed by cloud provider
- Makes it easier to handle biggest challenges identified on OnPrem version
- Costs are mainly
 - Computing from worker nodes
 - Control Plane HA
 - In some cases, additional services may have additional costs

Cloud or OnPrem



- It depends 😊
 - Data sovereignty
 - Security concerns (can be argued)
- Why not hybrid?
- Most clouds have services to add OnPrem Kubernetes Clusters to be managed as a Cloud Resource (Azure Arc, Amazon EKS Anywhere, Google Anthos)
- Allow to easily use additional services like Policies

CLI Tooling



CLI Tooling



- Preferable way to interact with Kubernetes Clusters is using CLI (Command-Line Interface) tools
- Motivation
 - Allow to automatize interactions
 - Needs much less resources
 - Small attach surface

API Server



- API Server is a “simple” REST API application, you can manage a cluster by making REST calls to the API Server
- Example HTTP Request: **GET /api/v1/namespaces/default/pods/{name}**
- Some challenges
 - Manage REST requests
 - Authn/authz
- Kubernetes have a CLI tool to make this interaction easily: **kubectl**

Kubectl



- API Server is a “simple” REST API application, you can manage a cluster by making REST calls to the API Server
- Example HTTP Request: **GET /api/v1/namespaces/default/pods/{name}**
- Some challenges
 - Manage REST requests
 - Authn/authz
- Kubernetes have a CLI tool to make this interaction easily: **kubectl**
- First challenge, how to pronounce? [‘kubectl’: The definitive pronunciation guide](#)

Kubectl: How to use?



`kubectl [command] [type] [name] [flags]`

- **command**: Operation that you want to perform on one or more resources, for example **create**, **get**, **describe**, **delete**.
- **type**: Resource type. Case-insensitive and can specify the singular, plural, or abbreviated forms.
- **name**: Case-insensitive name of the resource. If the name is omitted, details for all resources are displayed
- **flags**: Optional flags. For example, **-o** allow to specify output type of the commands

Kubectl: Examples



- List all nodes

```
kubectl get nodes
```

- Get more details on node node01

```
kubectl describe node node01
```

- List all nodes with more details

```
kubectl get node -o wide
```

- List all nodes and labels

```
kubectl get node --show-labels
```



Kubectl: How to reference resources?

- You can perform an action on several resource using only one command even on resources from different types
- Resources from same type

```
kubectl get pod pod-01 pod-02
```

- Resources from different types

```
kubectl get pod/pod-01 node/node01
```

[kubectl Cheat Sheet | Kubernetes](#)

[Kubectl Reference Docs \(kubernetes.io\)](#)

Kubectl: How to identify?



- **kubeconfig** is a file used to organize access to several cluster usually stored at **~/.kube/config**
- Needs to be kept on a secure place since have complete information about authn/authz of a user to a cluster
- This file should never be included on a repo or used for CI/CD process due to security reasons
- Uses the concept of **contexts**. A context makes a match between a cluster (API Server URL) and a user

Kubectl: Examples



- Get all clusters configuration available

```
kubectl config view
```

- Get actual context

```
kubectl config current-context
```

- Set another context

```
kubectl config use-context my-cluster-name
```



Demo | kubectl





Krew: Plugin Manager for Kubectl

- Allow to search, install and update additional tools to use with kubectl
- 210 plugins available: <https://krew.sigs.k8s.io/plugins/>
- Maintenance done by a Kubernetes SIG (Special Interest Group)
- List installed plugins

```
kubectl krew list
```

- Install a plugin

```
kubectl krew install <package>
```

Krew: Plugin Manager for Kubectl



- [cert-manager](#): Manage cert-manager resources inside your cluster
- [ctx](#): Switch between contexts in your kubeconfig
- [ingress-nginx](#): Interact with ingress-nginx
- [kubeseccan](#): Scan Kubernetes resources with kubeseccan.io.
- [ns](#): Switch between Kubernetes namespaces
- [resource-capacity](#): Provides an overview of resource requests, limits, and utilization

Declarative Configuration



Declarative Configuration



- An **imperative configuration** explicitly instructs a system on the steps to take to achieve a desired outcome (like using Docker commands):
 - Connect to container registry
 - Pull desired image
 - Create container
 - Start container
- A **declarative configuration** specifies a final, or desired state of an object, and lets the system determine what steps to take to achieve that state.
- The Kubernetes control plane continually and actively manages every object's **actual state to match the desired state** you supplied.

YAML: Yet Another Markup Language



- REST API applications like Kubernetes API Server, exchange data using JSON format
- When using kubectl, you provide a desired state configuration using the YAML markup language (Yet Another Markup Language)
- kubectl converts your YAML to JSON when communicate with the Kubernetes API Server
- When you create an YAML file to define Kubernetes resources, you're creating a [Kubernetes Manifest File](#)

YAML: Yet Another Markup Language



- Whitespace indentation is used to create file structure
- Tab characters are not allowed as part of that indentation
- Comments begin with the number sign (#) until the end of the line
- List members are denoted by a leading hyphen (-)
- An associative array entry is represented using colon space in the form key: value with one entry per line.
- Strings are ordinarily unquoted but may be enclosed in double-quotes ("), or single-quotes (').
- Multiple documents with single streams are separated with 3 hyphens (---).



Manifest File



- `apiVersion`: API group and version of the API you're calling to create this object
- `kind`: Object you want to create
- `metadata`: Data that helps uniquely identify the object, including a name string, UID, and optional namespace
- `spec` (most objects): Desired state for the object

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
spec:
  containers:
  - image: nginx:1.18-alpine
    name: nginx-container
    ports:
    - containerPort: 80
```

Pods



What is a Pod?



- Pods are the smallest deployable compute units you can create and manage in Kubernetes
- A Pod can manage one or more containers, with shared storage (volumes), environment variables, network resources, and a specification for how to run the containers
- Containers running in a Pod share the same IP and ports and communicate using native inter-process communication channels or localhost.
- Pods are immutable - if any change is made to the Pod specification (spec), a new Pod is created and then the old Pod is deleted

Pod Lifecycle



- Pod specification on a YAML file used by kubectl to ask the cluster to schedule the pod
- API Server add configuration in ETCD on a persistent way
- Scheduler finds a new pod maps to best available node
- Kubelet (on worker node) gets a notification about provisioning the pod and starts to create the associated containers
- Docker (or container runtime) creates new instances
- All pod status are saved on ETCD

Pod with one container

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
spec:
  containers:
  - image: nginx
    name: nginx-container
    resources: {}
    ports:
    - containerPort: 80
```

pod
name: nginx-pod

Container


Name: nginx-container

Port 80

Image: nginx

Pod with two containers

```
apiVersion: v1
kind: Pod
metadata:
  name: redis-nginx
  labels:
    app: web
spec:
  containers:
    - name: key-value-store
      image: redis
      ports:
        - containerPort: 6379
    - name: frontend
      image: nginx
      ports:
        - containerPort: 80
```



Kubectl: Work with pods



- Get all pods

```
kubectl get pods
```

- Get additional details from pods

```
kubectl get pods -o wide
```

- Get pod named **nginx-pod**

```
kubectl get pod nginx-pod
```

- Inspect pod named **nginx-pod**

```
kubectl describe pod nginx-pod
```

Kubectl: Interact with pods



- Get access to **nginx-pod** pod
`kubectl -it exec nginx-pod -- sh`
- Get access to container frontend on **redis-nginx** pod
`kubectl exec -it redis-nginx -c frontend -- bash`
- Port forwarding to port 80 on **nginx-pod** pod
`kubectl port-forward nginx-pod 8080:80`
- Port forwarding to port 80 on **redis-nginx** pod
`kubectl port-forward redis-nginx 8080:80`

Pods: Handle Resources



- A good practice when deploying pods on Kubernetes is to define the resources that will be used by it
- Kubernetes uses 2 concepts: Requests and Limits
- **Requests**: Amount of resources used by scheduler to define which node could best fit. This amount is always reserved for the pod
- **Limits**: Maximum amount of resources a pod can use.
- You can define request and limits for CPUs, memory and GPUs

Pods: Self Healing



- A Pod don't have any ability of self-healing
- Pods can be used directly but usually a controller is used to automatically manage your pods
- Each controller have specific way to control and manage their pods
 - ReplicaSets: Controls pods number of replicas
 - DaemonSets: Controls if one pods runs on each worker node
 - StatfulSets: Controls link between pod and persistent storage to handle pod state

Demo | Pods



Namespaces



Namespaces



- Kubernetes supports multiple virtual clusters backed by the same physical cluster
- These virtual clusters are called **namespaces** and creates logical partitions inside your cluster
- Namespaces provide a scope for names
- Names of resources need to be unique within a namespace, but not across namespaces.
- Namespaces can be used to
 - Allowing/denying network communications
 - Define RBAC permissions
 - Define resource quotas

Namespaces are mandatory?



- Usually, your resources are created in the context of a namespace
- “**default**” namespace is created on a empty cluster
- Your **kubectl** commands uses the namespace you set on your **kubeconfig** to be the default
- You can use **kubectl ns** plugin to easily change it
- When you want to not change the default, you need to use the flag **-n** in each command



Namespaces are mandatory?



- When a resource is not namespaced, means that is a cluster resource
- Get list of resources that aren't namespaced scope

```
kubectl api-resources --namespaced=false
```

NAME	SHORTNAMES	APIVERSION	NAMESPACED	KIND
componentstatuses	cs	v1	false	ComponentStatus
namespaces	ns	v1	false	Namespace
nodes	no	v1	false	Node
persistentvolumes	pv	v1	false	PersistentVolume
mutatingwebhookconfigurations		admissionregistration.k8s.io/v1	false	MutatingWebhookConfiguration
validatingwebhookconfigurations		admissionregistration.k8s.io/v1	false	ValidatingWebhookConfiguration
customresourcedefinitions	crd, crds	apiextensions.k8s.io/v1	false	CustomResourceDefinition

Kubectl: Work with namespaces



- Get all namespaces

```
kubectl get namespaces
```

- Get all namespaces using shortname

```
kubectl get ns
```

- Get list of pods from namespace **my-ns**

```
kubectl get pod -n my-ns
```

- Delete **my-ns** namespace

```
kubectl delete ns my-ns
```

Namespaces: Why to use?



- Namespaces can be used to define global policies
 - Allowing/denying network communications
 - Define RBAC permissions
 - Define resource quotas
- Namespaces are usually created by:
 - Component Type – Ex: All backends in one namespace, all websites in another
 - Users – Ex: User rights/quotas can be limited by namespace
 - Environments – Ex: Dev resources can be in one namespace, QA in another
 - System Segment – Ex: Catalog microservices in one namespace, ordering in another

Namespaces: How to define?



- Namespaces are usually created by:
 - Component Type – Ex: All backends in one namespace, all websites in another
 - Users – Ex: User rights/quotas can be limited by namespace
 - Environments – Ex: Dev resources can be in one namespace, QA in another
 - System Segment – Ex: Catalog microservices in one namespace, ordering in another

Demo | Namespaces



Questions?



kubernetes

Lab #01: Introduction to K8S



kubernetes